

Lakshyaraj Singh Rao

****Full-Stack Engineer • AI Systems • Backend • DevOps****

Mumbai, India • raolakshyaraj@gmail.com • [linkedin.com/in/lakshyaraj](https://linkedin.com/in/laks

SUMMARY

Backend-focused engineer who ships production AI products end-to-end. Solo-built and operates Homest

TECHNICAL SKILLS

****Languages:**** TypeScript, JavaScript (ES6+), C++, SQL

****Backend:**** Node.js, Express.js, Next.js 15, REST APIs, JWT Auth, Prisma ORM, Zod

****AI / LLM:**** GPT-4o, Claude Sonnet, Vercel AI SDK, RAG (pgvector), Streaming, Prompt Engineering

****Databases:**** PostgreSQL (Neon), pgvector, MongoDB, Redis (Upstash), RabbitMQ

****DevOps & Cloud:**** Docker, Docker Compose, GitHub Actions, Vercel, AWS EC2, Kubernetes, ArgoCD

****Observability:**** Sentry, Prometheus, Grafana, Winston, Playwright, Vitest

EXPERIENCE

Founding Engineer & Co-Founder â Homesty AI Technology LLP

*Mumbai, India • Mar 2026 â Present • Entity incorporating May 2026 • [homesty.ai](https://homesty.

Solo engineer on the production stack. End-to-end ownership: system prompt design, RAG infrastruc

- Shipped buyerchat to production: Next.js 15 + Postgres + pgvector + Prisma 7 + GPT-4o + Claude +
- 165 production deployments. 0 critical Sentry classes firing under live traffic. 159 unit tests
- Engineered 5-layer anti-fabrication architecture: PART 0 hard-stop rules + onChunk markdown abor
- Closed 8 production hallucination classes including OTP-simulation, fabricated entity names, and
- 8-stage chat pipeline: per-IP rate limit â intent classifier â RAG retrieval â GPT-4o streaming
- Built 7-module Decision Engine: scores projects against buyer preferences, generates Strong Buy
- pgvector retriever with 600ms timeout + 0.30 cosine floor (silently drops weak matches instead
- Security: prompt injection protection, PII stripping, origin checks, PRICING_LOCKED canonical su

PROJECTS

[rag-starter](https://github.com/ykstorm/rag-starter)

Production RAG template extracted from buyerchat. pgvector retrieval, embed-writer, cosine-floor, pr

[guardrail-proxy](https://github.com/ykstorm/guardrail-proxy)

Streaming LLM safety. Mid-chunk abort when audit catches violation â user never sees a bad response

[devops-showcase](https://github.com/ykstorm/devops-showcase)

Kubernetes + ArgoCD GitOps stack. Canary deploys, Prometheus, Grafana. Apache 2.0.

[k8s-deploy-tracker](https://github.com/ykstorm/k8s-deploy-tracker)

Go service that consumes Kubernetes deploy events and feeds failing ones to GPT-4 for one-paragraph

[stream-bench](https://github.com/ykstorm/stream-bench)

Benchmarks for LLM streaming: time-to-first-token, throughput, total latency across OpenAI, Anthropi

EDUCATION

****Manipal University Jaipur**** â B.Tech, Computer Science

Graduating 2026

RECOGNITION

- ****Production Reliability:**** 165 deployments shipped, 0 critical Sentry classes firing, 159 unit
- ****Anti-Fabrication Architecture:**** Closed 8 production hallucination classes via PART 0 hard rul