

Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Full-Stack · DevOps

Mumbai, India · Open to Bangalore / Remote
raolakshyaraj@gmail.com
lakshyaraj-dev.vercel.app · GitHub · LinkedIn · npm

SUMMARY

I build the reliability layer that lets AI systems run in production without supervision — refusal when retrieval fails, abort when output goes off-script, eval-gating before merge, idempotency under webhook re-delivery. Co-Founder and sole engineer on Homesty.ai LLP since Nov 2025, serving real traffic on Next.js 15 + pgvector + GPT-4o. Four @ykstorsorg npm packages (one with SLSA provenance) and a public GitHub Marketplace Action — extracted from what production actually broke.

TECHNICAL SKILLS

Languages

TypeScript · JavaScript
Python · SQL
Rust (learning)

Backend / Data

Node.js · Express · Next.js
PostgreSQL · pgvector
Prisma 7 · Redis · BullMQ

Infra / AI

Docker · Kubernetes · Helm
Terraform · GitHub Actions · Vercel
OpenAI · Anthropic · Ollama

OPEN-SOURCE PROJECTS

Anchor [Creator](#) · anchor-iota-ten.vercel.app/playground

- Every RAG demo breaks when retrieval finds nothing useful — the AI just makes things up. **Anchor is the RAG layer that refuses instead of guessing: above the similarity floor it returns chunks with provenance; below, it returns refused: true.** Built into a real-estate production app where a hallucination is a lawsuit. Live playground, one-command Docker setup.

Postgres + pgvector Prisma 7 Next.js 15 OpenAI embeddings

Codecraft [Creator](#) · codecraft-ai-tau.vercel.app

- “Try-it” links die when users have to install Node, clone the repo, and run `npm install`. **Codecraft runs a real Vite + React dev server inside the browser tab — edit code in Monaco, type `npm install dayjs` in the terminal, watch the preview hot-reload in under 2 seconds.** No backend, no install; the COOP/COEP cross-origin-isolation boundary unlocks `SharedArrayBuffer`, and an IndexedDB snapshot cache drops repeat-visit boots under 20s.

WebContainers Monaco xterm Vite COOP/COEP

Anvil [Creator](#) · github.com/ykstorm/anvil

- Stripe sends the same webhook three times when the internet hiccups, and your worker double-charges users. **Anvil is the dedup + retry + dead-letter pipeline that handles every webhook source (Stripe, GitHub, Twilio) so your worker never repeats.** Drop-in library, `\u{40}ykstorsorg/anvil v0.1.1` on npm with SLSA build provenance. Ships a Terraform module and a Helm chart. Five design contracts proven by tests in CI against real Redis.

TypeScript BullMQ Redis Express Terraform Helm

Tripwire [Creator](#) · github.com/ykstorm/tripwire

- Once an LLM token reaches the user it's already wrong — post-hoc moderation runs after the damage. **Tripwire watches each token mid-stream and kills the response on rule trip: PII leaks, prompt-injection, fabricated facts.** Drop in as an OpenAI-compatible proxy (change one URL) or import as a library. Sub-millisecond per chunk; zero buffering on the happy path. `\u{40}ykstorsorg/tripwire v1.1.0` on npm + Docker image at `ghcr.io/ykstorm/tripwire`.

TypeScript Node async iterators streaming abort Docker

Goldset [Creator](#) · [Maintainer](#) · github.com/ykstorm/goldset

- “Did my prompt change break anything?” — today you eyeball outputs in a notebook and ship a guess. **Goldset is CI for AI: three eval runners (golden answers, LLM judge, structural assertions) run on every PR, post a delta-vs-base comment, and block merge on regression.** Pluggable OpenAI / Anthropic judge. `\u{40}ykstorsorg/goldset v0.2.4` on npm + GitHub Marketplace.

TypeScript esbuild OpenAI Anthropic GitHub Actions

Quickdraw [Creator](#) · github.com/ykstorm/quickdraw

- Every AI provider claims to be “fastest” — but fast on what prompts, what network, what time of day? **Quickdraw measures real LLM streaming on your prompts: TTFT, tokens/sec, p50/p95/p99, and cost per 1K across OpenAI + Anthropic.** `quickdraw diff` catches regressions between nightly runs. `\u{40}ykstorsorg/quickdraw v1.0.3` on npm with SLSA build provenance.

TypeScript OpenAI Anthropic CLI

Stackup [Creator](#) · github.com/ykstorm/stackup

- Production Kubernetes patterns cost \$200+/month on cloud, and most “K8s tutorials” stop at `kubectl run nginx`. **Stackup runs the full production stack on your laptop in 15 minutes: ArgoCD GitOps app-of-apps, Argo Rollouts canary with real Prometheus success-rate analysis (verified live: 25 – 50 – 75 – 100% with the ≥ 0.95 gate passing 3 $\times$), kube-prometheus-stack monitoring.** Real progressive delivery, free.

Kubernetes ArgoCD Helm kind Grafana

EXPERIENCE

Homesty.ai LLP [Co-Founder & Sole Engineer](#) · Remote · Nov 2025 – Present

- Sole engineer on a live buyer-side real-estate AI: Next.js 15, Postgres + pgvector, Prisma, GPT-4o and Claude streaming, behind rate limiting, intent routing, and Sentry.
- Built the chat path around refusal-first retrieval and a mid-stream guardrail — the AI declines when it has no grounded source rather than improvising, which in real estate is a liability, not a UX choice. Both patterns were hardened into standalone OSS (Anchor, Tripwire).

Next.js 15 Postgres + pgvector Prisma 7 GPT-4o Claude Sentry Docker

EDUCATION

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026
Coursework: Distributed Systems · Database Internals · Networks · Machine Learning Foundations.
Devagiri CMI Public School, Calicut — CBSE Class XII, 84%
Vidya Kendra, Calicut — CBSE Class X, 94%