

Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Full-Stack · DevOps

Mumbai, India · Open to Bangalore / Remote · [homesty.ai](#) · [raolakshyaraj@gmail.com](#) · [lakshyaraj-dev.vercel.app](#) · [GitHub](#) · [LinkedIn](#) · [npm](#)

SUMMARY

I build the reliability layer that lets AI systems run in production without supervision — refusal when retrieval fails, abort when output goes off-script, evaluating before merge, idempotency under webhook re-delivery. Co-Founder and sole engineer on Homesty.ai LLP since Nov 2025, serving real traffic on Next.js 15 + pgvector + GPT-4o. Four [@ykstorsorg](#) npm packages (one with SLSA provenance) and a public GitHub Marketplace Action — extracted from what production actually broke.

EXPERIENCE

Homesty.ai LLP · [Co-Founder & Founding Engineer](#) · Remote · Nov 2025 – Present

- Sole engineer on a live buyer-side real-estate AI: Next.js 15, Postgres + pgvector, Prisma, GPT-4o and Claude streaming, behind rate limiting, intent routing, and Sentry. Built the chat path around refusal-first retrieval and a mid-stream guardrail — the AI declines when it has no grounded source rather than improvising, which in real estate is a liability, not a UX choice. Both patterns were hardened into standalone OSS (Anchor, Tripwire).

[Next.js 15](#) [Postgres + pgvector](#) [Prisma 7](#) [GPT-4o](#) [Claude](#) [Sentry](#) [Docker](#)

PROJECTS

Anchor · [Creator](#) · [anchor-iota-ten.vercel.app/playground](#)

- Every RAG demo breaks when retrieval finds nothing useful — the AI just makes things up. **Anchor is the RAG layer that refuses instead of guessing: above the similarity floor it returns chunks with provenance; below, it returns refused: true.** Built into a real-estate production app where a hallucination is a lawsuit. Live playground; p50 2.6 ms retrieval at 100k vectors (HNSW), measured in CI.

[Postgres + pgvector](#) [Prisma 7](#) [Next.js 15](#) [OpenAI embeddings](#)

Anvil · [Creator](#) · [github.com/ykstorm/anvil](#)

- Stripe sends the same webhook three times when the internet hiccups, and your worker double-charges users. **Anvil is the dedup + retry + dead-letter pipeline that handles every webhook source (Stripe, GitHub, Twilio) so your worker never repeats.** [\u{40}ykstorsorg/anvil v0.1.1](#) on npm with SLSA provenance; Terraform module + Helm chart. 10.5k req/s ingress, constant-time HMAC verify (0.9% timing delta) — measured in CI.

[TypeScript](#) [BullMQ](#) [Redis](#) [Express](#) [Terraform](#) [Helm](#)

Tripwire · [Creator](#) · [github.com/ykstorm/tripwire](#)

- Once an LLM token reaches the user it's already wrong — post-hoc moderation runs after the damage. **Tripwire watches each token mid-stream and kills the response on rule trip: PII leaks, prompt-injection, fabricated facts.** Drop in as an OpenAI-compatible proxy (one URL change) or a library — 4.7 μ s per chunk (measured in CI), zero buffering on the happy path. [\u{40}ykstorsorg/tripwire v1.1.0](#) on npm + Docker at [ghcr.io/ykstorm/tripwire](#).

[TypeScript](#) [Node async iterators](#) [streaming abort](#) [Docker](#)

Goldset · [Creator](#) · [Maintainer](#) · [github.com/ykstorm/goldset](#)

- “Did my prompt change break anything?” — today you eyeball outputs in a notebook and ship a guess. **Goldset is CI for AI: three eval runners (golden answers, LLM judge, structural assertions) run on every PR, post a delta-vs-base comment, and block merge on regression.** [\u{40}ykstorsorg/goldset v0.2.4](#) on npm + GitHub Marketplace; deterministic runners clear 1k cases in ~3 ms, so the gate never blocks a PR.

[TypeScript](#) [esbuild](#) [OpenAI](#) [Anthropic](#) [GitHub Actions](#)

Stackup · [Creator](#) · [github.com/ykstorm/stackup](#)

- Production Kubernetes patterns cost \$200+/month on cloud, and most “K8s tutorials” stop at `kubectl run nginx`. **Stackup runs the full production stack on your laptop in 15 minutes: ArgoCD GitOps app-of-apps, Argo Rollouts canary with real Prometheus success-rate analysis, kube-prometheus-stack monitoring.** Canary CI-verified: kind + Argo Rollouts advancing 25 → 50 → 75 → 100% only after a live Prometheus ≥ 0.95 success-rate gate passes.

[Kubernetes](#) [ArgoCD](#) [Helm](#) [kind](#) [Grafana](#)

Quickdraw · [Creator](#) · [github.com/ykstorm/quickdraw](#)

- Every AI provider claims to be “fastest” — but fast on what prompts, what network, what time of day? **Quickdraw measures real LLM streaming on your prompts: TTFT, tokens/sec, p50/p95/p99, and cost per 1K across OpenAI + Anthropic.** [quickdraw diff](#) catches regressions between nightly runs. [\u{40}ykstorsorg/quickdraw v1.0.3](#) on npm with SLSA build provenance.

[TypeScript](#) [OpenAI](#) [Anthropic](#) [CLI](#)

TECHNICAL SKILLS

Languages

TypeScript · JavaScript
Python · SQL
Rust (learning)

Backend / Data

Node.js · Express · Next.js
PostgreSQL · pgvector
Prisma 7 · Redis · BullMQ

Infra / AI

Docker · Kubernetes · Helm
Terraform · GitHub Actions · Vercel
OpenAI · Anthropic · Ollama

EDUCATION

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026

Coursework: Distributed Systems · Database Internals · Networks · Machine Learning Foundations.

Devagiri CMI Public School, Calicut — CBSE Class XII, 84%

Vidya Kendra, Calicut — CBSE Class X, 94%