

Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Eval Systems

Mumbai → Remote · balveer767@gmail.com · lakshyaraj-dev.vercel.app · github.com/ykstorm · linkedin.com/in/lakshyaraj

Summary

Backend engineer specializing in production AI infrastructure: provenance-first RAG (Anchor), mid-stream LLM guardrails (Tripwire), and CI-blocking eval harnesses (Goldset). Author of three npm packages under @ykstormsorg and a public GitHub Action listed on the GitHub Marketplace. Apache 2.0 by default.

Technical Skills

Languages

TypeScript, Python,
Rust (learning), SQL

Backend / Data

Node.js, Postgres,
pgvector, Prisma 7,
BullMQ, Redis

Infra / DevOps / AI

Docker, Kubernetes,
Terraform, GitHub Actions,
OpenAI, Anthropic, Ollama

Open-Source Projects

Goldset — Three-runner AI eval framework (goldenDataset · llmJudge · structural) shipped as a public GitHub Action. PR comment diffs the delta; structural-assertion failures block merge. 938 KB self-contained esbuild bundle — zero install at consumer side.

Published: @ykstormsorg/goldset v0.2.0 · GitHub Marketplace (Continuous Integration)

TypeScript · esbuild · OpenAI · Anthropic · Apache 2.0

Anchor — Provenance-first RAG with cosine-floor refusal. Returns chunks when similarity ≥ 0.30 ; returns refused: true and zero chunks below floor. Idempotent upsert keyed on (sourceType, sourceId). 60 embeddings across 16 RERA-verified projects in demo corpus.

Live: anchor-iota-ten.vercel.app/playground

TypeScript · Postgres + pgvector · Prisma 7 · Next.js 15 · OpenAI text-embedding-3-small

Tripwire — Mid-stream LLM safety. Token-by-token rule engine aborts the response stream on rule trip — user never sees a bad response finish rendering. Sub-millisecond decision per chunk, zero buffering on happy path.

Published: @ykstormsorg/tripwire

TypeScript · Node async iterators · streaming abort

Quickdraw — LLM streaming benchmark CLI capturing TTFT, tokens-per-second, and dollar-per-1K across OpenAI and Anthropic. Nightly run at 03:00 UTC diffs vs. baseline; results auto-committed to repo.

Published: @ykstormsorg/quickdraw

TypeScript · OpenAI SDK · Anthropic SDK · GitHub Actions cron

Stackup — Local Kubernetes lab spinning up via single make up in under 10 minutes. ArgoCD + Argo Rollouts + Grafana + Prometheus on kind cluster — production-shaped stack for learning, not for HA workloads.

Code: github.com/ykstorm/stackup

Helm · ArgoCD · kind · Bash

Codecraft — In-browser IDE prototype. WebContainers running real Node.js + Monaco editor + Ollama for local AI completions. Landing + auth shell deployed; IDE workspace in active build.

Live: codecraft-ai.vercel.app

WebContainers · Monaco · xterm · Next.js

In Progress

Anvil — Backend Platform Module Suite — Idempotent BullMQ job queue + retry-safe webhook receiver with exponential backoff and dead-letter handling + Terraform module (Postgres + Redis + worker pool) + Helm chart. Targeting backend-platform roles at Supabase, Trigger.dev, Razorpay-tier.

Experience

Freelance Backend & AI Engineer — Remote · Mar 2025 – Present

Production AI features for client projects: RAG pipelines (pgvector + Prisma), LLM streaming with guardrails, CI eval harnesses, Docker + GitHub Actions deploy pipelines. Multiple deployments live; client specifics under NDA.

Education

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026

Coursework: Distributed Systems, Database Internals, ML Foundations, Networks

Recognition

- 3 npm packages published under @ykstormsorg scope: goldset, quickdraw, tripwire
- Goldset listed on GitHub Marketplace (Continuous Integration category)
- 6 open-source repositories under Apache 2.0 with working demos
- Portfolio site with live npm + GitHub activity at lakshyaraj-dev.vercel.app

Latest version: github.com/ykstorm/resume