Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Full-Stack · DevOps

Mumbai, India · Open to Bangalore / Remote · homesty.ai · raolakshyaraj@gmail.com · lakshyaraj-dev.vercel.app · GitHub · LinkedIn · npm

SUMMARY

I build the reliability layer that lets AI systems run in production without supervision — refusal when retrieval fails, abort when output goes off-script, eval-gating before merge, idempotency under webhook re-delivery. Co-Founder and sole engineer on Homesty.ai LLP since Nov 2025, serving real traffic on Next.js 15 + pgvector + GPT-4o. Four @ykstormsorg npm packages (one with SLSA provenance) and a public GitHub Marketplace Action — extracted from what production actually broke.

EXPERIENCE

Homesty.ai LLP Co-Founder & Founding Engineer · Remote · Nov 2025 – Present

- Sole engineer on a live buyer-side real-estate AI: Next.js 15, Postgres + pgvector, Prisma, GPT-4o and Claude streaming, behind rate limiting, intent routing, and Sentry. Built the chat path around refusal-first retrieval and a mid-stream guardrail — the AI declines when it has no grounded source rather than improvising, which in real estate is a liability, not a UX choice. Both patterns were hardened into standalone OSS (Anchor, Tripwire).

`Next.js 15` `Postgres + pgvector` `Prisma 7` `GPT-4o` `Claude` `Sentry` `Docker`

PROJECTS

Anchor Creator · anchor-iota-ten.vercel.app/playground

- Every RAG demo breaks when retrieval finds nothing useful — the AI just makes things up. **Anchor is the RAG layer that refuses instead of guessing: above the similarity floor it returns chunks with provenance; below, it returns `refused: true`.** Built into a real-estate production app where a hallucination is a lawsuit. Live playground; p50 2.6 ms retrieval at 100k vectors (HNSW), measured in CI.

`Postgres + pgvector` `Prisma 7` `Next.js 15` `OpenAI embeddings`

Anvil Creator · github.com/ykstorm/anvil

- Stripe sends the same webhook three times when the internet hiccups, and your worker double-charges users. **Anvil is the dedup + retry + dead-letter pipeline that handles every webhook source (Stripe, GitHub, Twilio) so your worker never repeats.** `\u{40}ykstormsorg/anvil` v0.1.1 on npm with SLSA provenance; Terraform module + Helm chart. 10.5k req/s ingress, constant-time HMAC verify (0.9% timing delta) — measured in CI.

`TypeScript` `BullMQ` `Redis` `Express` `Terraform` `Helm`

Tripwire Creator · github.com/ykstorm/tripwire

- Once an LLM token reaches the user it's already wrong — post-hoc moderation runs after the damage. **Tripwire watches each token mid-stream and kills the response on rule trip: PII leaks, prompt-injection, fabricated facts.** Drop in as an OpenAI-compatible proxy (one URL change) or a library — 4.7 µs per chunk (measured in CI), zero buffering on the happy path. `\u{40}ykstormsorg/tripwire` v1.1.0 on npm + Docker at `ghcr.io/ykstorm/tripwire`.

`TypeScript` `Node async iterators` `streaming abort` `Docker`

Goldset Creator · Maintainer · github.com/ykstorm/goldset

- "Did my prompt change break anything?" — today you eyeball outputs in a notebook and ship a guess. **Goldset is CI for AI: three eval runners (golden answers, LLM judge, structural assertions) run on every PR, post a delta-vs-base comment, and block merge on regression.** `\u{40}ykstormsorg/goldset` v0.2.4 on npm + GitHub Marketplace; deterministic runners clear 1k cases in 3 ms, so the gate never blocks a PR.

`TypeScript` `esbuild` `OpenAI` `Anthropic` `GitHub Actions`

Stackup Creator · github.com/ykstorm/stackup

- Production Kubernetes patterns cost $200+/month on cloud, and most "K8s tutorials" stop at `kubectl run nginx`. **Stackup runs the full production stack on your laptop in 15 minutes: ArgoCD GitOps app-of-apps, Argo Rollouts canary with real Prometheus success-rate analysis, kube-prometheus-stack monitoring.** Canary verified live: 25→50→75→100% with the ≥0.95 analysis gate passing 3×.

`Kubernetes` `ArgoCD` `Helm` `kind` `Grafana`

Quickdraw Creator · github.com/ykstorm/quickdraw

- Every AI provider claims to be "fastest" — but fast on what prompts, what network, what time of day? **Quickdraw measures real LLM streaming on your prompts: TTFT, tokens/sec, p50/p95/p99, and cost per 1K across OpenAI + Anthropic.** `quickdraw diff` catches regressions between nightly runs. `\u{40}ykstormsorg/quickdraw` v1.0.3 on npm with SLSA build provenance.

`TypeScript` `OpenAI` `Anthropic` `CLI`

TECHNICAL SKILLS

Languages	Backend / Data	Infra / AI
TypeScript · JavaScript	Node.js · Express · Next.js	Docker · Kubernetes · Helm
Python · SQL	PostgreSQL · pgvector	Terraform · GitHub Actions · Vercel
Rust (learning)	Prisma 7 · Redis · BullMQ	OpenAI · Anthropic · Ollama

EDUCATION

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026
Coursework: Distributed Systems · Database Internals · Networks · Machine Learning Foundations.
Devagiri CMI Public School, Calicut — CBSE Class XII, 84%
Vidya Kendra, Calicut — CBSE Class X, 94%