

Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Eval Systems

Mumbai, India · [raolakshyaraj at gmail.com](mailto:raolakshyaraj@gmail.com) · lakshyaraj-dev.vercel.app · github.com/ykstorm

Summary

Backend engineer specializing in production AI infrastructure: provenance-first RAG, mid-stream LLM guardrails, and eval harnesses that block bad merges. Author of three published npm packages under the @ykstormsorg scope and a public GitHub Action that posts PR comments and blocks merges on AI behavior regression.

Open-Source Projects

Goldset — Three-runner AI eval framework (goldenDataset, llmJudge, structural) shipped as a public GitHub Action. PR comment diffs the delta; structural-assertion failures block merge. Single bundled dist/action.cjs (938 KB) — no install at consumer side. github.com/ykstorm/goldset · npm

TypeScript · esbuild · OpenAI · Anthropic · Apache 2.0

Anchor — Provenance-first RAG layer with cosine-floor refusal. Returns chunks when similarity ≥ 0.30 ; returns refused: true and zero chunks below floor. Idempotent upsert on (sourceType, sourceId). anchor-iota-ten.vercel.app

TypeScript · Next.js 15 · Postgres + pgvector · Prisma 7 · OpenAI text-embedding-3-small

Tripwire — Mid-stream LLM safety. Token-by-token rule engine aborts the response stream on rule trip. Sub-millisecond decision per chunk, zero buffering on happy path. github.com/ykstorm/tripwire

TypeScript · Node async iterators · streaming abort

Quickdraw — LLM streaming benchmark CLI. Captures TTFT, tokens-per-second, and dollar-per-1K across OpenAI and Anthropic. Nightly run at 03:00 UTC diffs vs baseline. github.com/ykstorm/quickdraw

TypeScript · OpenAI SDK · Anthropic SDK

Stackup — Production-style Kubernetes locally in 10 minutes via single make up target. ArgoCD + Argo Rollouts + Grafana + Prometheus on kind cluster. github.com/ykstorm/stackup

Helm · ArgoCD · kind · Bash

Codecraft — In-browser IDE running real Node.js via WebContainers, Monaco for editing, local Ollama for AI completions. Zero cloud roundtrip required. codecraft-ai.vercel.app

WebContainers · Monaco · Ollama · Next.js

In Progress

Backend Platform Module Suite — BullMQ-backed idempotent job queue + retry-safe webhook receiver + Terraform module for one-click stack provisioning. Targeting backend-platform roles at Supabase, Trigger.dev, Razorpay.

Experience

Freelance Backend & AI Engineer — Remote · Mar 2025 – Present Production AI features for client projects: RAG pipelines, LLM streaming, eval harnesses, and CI/CD pipelines. Multiple deployments live; client specifics under NDA.

Technical Skills

- **Languages:** TypeScript, Python, Rust (learning)
- **Backend:** Node.js, Postgres, pgvector, Prisma 7, BullMQ, Redis
- **Infra:** Docker, Kubernetes, Terraform, GitHub Actions, Vercel
- **AI / Retrieval:** OpenAI, Anthropic, pgvector, Ollama, eval harnesses, RAG

Education

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026

Recognition

- 3 npm packages published under the @ykstormsorg scope
- Goldset listed on GitHub Marketplace (Continuous Integration)
- 6 open-source repos under Apache 2.0 with working demos
- Anchor live demo + playground at anchor-iota-ten.vercel.app

