

Lakshyaraj Singh Rao

Full-Stack Engineer · AI Systems · Backend · DevOps

Mumbai → Remote · balveer767@gmail.com · lakshyaraj-dev.vercel.app · github.com/ykstorm

Summary

Full-stack engineer specialized in AI infrastructure: vector retrieval, LLM streaming guardrails, eval harnesses. Built 6 OSS projects + 3 npm packages in the AI tooling space. Apache 2.0 by default.

Recent Projects

Goldset — Golden datasets + LLM-as-judge + structural assertions, as a GitHub Action. Three-runner architecture posts PR comment with delta, blocks merge on regression. *Published on npm + GitHub Marketplace.* github.com/ykstorm/goldset

Anchor — Provenance-first RAG with cosine-floor refusal. Postgres + pgvector + Prisma + Next.js. Live demo with playground. *16 projects, 5 builders, 60 embeddings.* anchor-iota-ten.vercel.app

Tripwire — Mid-stream LLM token guardrail. Sub-millisecond abort on rule trip. Token-by-token pattern matching against rule set. github.com/ykstorm/tripwire

Quickdraw — LLM streaming benchmark CLI. Captures TTFT, tokens/sec, USD per 1K tokens. Nightly bench against OpenAI + Anthropic, diffs vs baseline. github.com/ykstorm/quickdraw

Stackup — Local production-style Kubernetes in 10 minutes. ArgoCD + Argo Rollouts + Grafana via single make up target. github.com/ykstorm/stackup

Codecraft — In-browser IDE running real Node.js via WebContainers + Monaco + local Ollama for AI completions. github.com/ykstorm/codecraft-ai

Technical Skills

- **Languages:** TypeScript, Python, Rust (learning)
- **Backend:** Node.js, Postgres, Redis, Prisma, BullMQ
- **Infra:** Docker, Kubernetes, Terraform, GitHub Actions, Vercel
- **AI / Retrieval:** OpenAI, Anthropic, pgvector, Ollama, eval harnesses, RAG

Apache 2.0 by default. Latest version: github.com/ykstorm/resume