

Lakshyaraj Singh Rao

Backend Engineer · AI Infrastructure · Full-Stack · DevOps

Jaipur, India · Remote
raolakshyaraj@gmail.com
lakshyaraj-dev.vercel.app · GitHub · LinkedIn
· npm

SUMMARY

Backend-focused engineer building production AI infrastructure. Four @ykstorsorg npm packages (one with build provenance) and a public GitHub Marketplace Action. Patterns from a solo production AI deployment, extracted into seven open-source repos: RAG, LLM safety, eval, benchmarking, Kubernetes/GitOps, webhook reliability.

TECHNICAL SKILLS

Languages

TypeScript · JavaScript
Python · SQL
Rust (learning)

Backend / Data

Node.js · Express · Next.js
PostgreSQL · pgvector
Prisma 7 · Redis · BullMQ

Infra / AI

Docker · Kubernetes · Helm
Terraform · GitHub Actions · Vercel
OpenAI · Anthropic · Ollama

OPEN-SOURCE PROJECTS

Anvil [Creator](#) · [github.com/ykstorm/anvil](#)

- Idempotent webhook → BullMQ worker pipeline: HMAC-SHA256 constant-time verify, dedupe by sha256(signature + payload), retry backoff [1s, 5s, 30s, 5m], dead-letter replay via a separate consumer.
- Five design contracts proven by tests in CI (Redis-backed). Ships a Terraform module + Helm chart. Published as `ykstorsorg/anvil` v0.1.0 on npm with build provenance.

TypeScript BullMQ Redis Express Terraform Helm

Goldset [Creator](#) · [Maintainer](#) · [github.com/ykstorm/goldset](#)

- Three-runner AI eval framework (goldenDataset · llmJudge · structural assertions) shipped as a public GitHub Action.
- PR comment diffs the delta on every push; structural-assertion failures block merge. Published as `ykstorsorg/goldset` v0.2.0 on npm + GitHub Marketplace.

TypeScript esbuild OpenAI Anthropic GitHub Actions

Anchor [Creator](#) · [anchor-iota-ten.vercel.app/playground](#)

- Provenance-first RAG layer with cosine-floor refusal. Returns chunks when similarity ≥ 0.30 ; returns refused: true and zero chunks below floor.
- Live playground renders the retrieval decision — chunks above the floor, or an explicit refusal — for any query.
- Adaptive K per query intent, with a provenance array on every response so answers are traceable to source.

Postgres + pgvector Prisma 7 Next.js 15 OpenAI embeddings

Tripwire [Creator](#) · [github.com/ykstorm/tripwire](#)

- Mid-stream LLM safety. Token-by-token rule engine aborts the response stream on rule trip, before the bad token reaches the user.
- Sub-millisecond decision per chunk; zero buffering on happy path. Published as `ykstorsorg/tripwire` on npm.

TypeScript Node async iterators streaming abort

Quickdraw [Creator](#) · [github.com/ykstorm/quickdraw](#)

- LLM streaming benchmark CLI — TTFT, tokens/sec, cost per 1K tokens. Nightly run across OpenAI + Anthropic, diffs vs. baseline. Published as `ykstorsorg/quickdraw` on npm.

TypeScript OpenAI SDK Anthropic SDK GHA cron

Stackup [Creator](#) · [github.com/ykstorm/stackup](#)

- Production-shape Kubernetes on a laptop in 10 minutes (single make up): ArgoCD app-of-apps + Argo Rollouts canary + kube-prometheus-stack on a kind cluster.

Kubernetes ArgoCD Helm kind Grafana

Codecraft [Creator](#) · [codecraft-ai-tau.vercel.app](#)

- Boots a real Next.js dev server inside the browser tab via WebContainers — no backend, no local install. The whole cross-origin-isolation (COOP/COEP) boundary is what makes SharedArrayBuffer and the in-tab runtime work.

WebContainers Monaco xterm Next.js COOP/COEP

FREELANCE WORK

Homesty.ai [Solo Backend & AI Engineer](#) · Remote · Mar 2025 – Present

- Sole engineer on a live buyer-side real-estate AI: Next.js 15, Postgres + pgvector, Prisma, GPT-4o and Claude streaming, behind rate limiting, intent routing, and Sentry.
- Built the chat path around refusal-first retrieval and a mid-stream guardrail — the AI declines when it has no grounded source rather than improvising, which in real estate is a liability, not a UX choice. Both patterns were hardened into standalone OSS (Anchor, Tripwire).

Next.js 15 Postgres + pgvector Prisma 7 GPT-4o Claude Sentry Docker

EDUCATION

Manipal University Jaipur — B.Tech, Computer Science · Graduating 2026

